What Is Flo State?

Clinical trial teams spend 60% of their time on administrative work. Not because they want to. Because the systems they rely on were never designed to work together.

70% of organizations are experimenting with AI. Only 12% use it consistently. The gap is not a technology problem. It is a training problem, an integration problem, and a trust problem.

Flo State closes that gap.

It is the first AI concierge built specifically for clinical trial operations. Not a chatbot. Not a general-purpose tool you have to configure yourself. Flo State is a curated library of purpose-built AI capabilities that work inside Florence, and speaks to all the other tools your team already uses every day. Everyday AI means exactly what it sounds like: intelligence that shows up in your workflow on Monday morning, not in a demo that never ships.

Three category types with one goal: make it easy for everyone on your team to use proven AI to work, so sites-CROs-sponsors can easily scale the value derived from those agentic use cases every day, instead of having your clinops workforce wasting time tinkering in ChatGPT or having to tell investors that big digital transformation initiative was an implementation of an AI-generated meeting transcript tool.

Three Category Types

Agents: AI that does work for you.

Templates: Pre-configured AI project workspaces for your role (i.e Claude, ChatGPT, Gemini, etc.).

Skills: Modular capabilities you can add anywhere.

Who Flo State Is For

If you work in clinical trial operations at a site, CRO, or sponsor organization, Flo State was built for your daily work. Specifically:

- Coordinators who juggle 7+ studies and spend more time on paperwork than patients.
- CRAs who need to review site documents faster without missing findings.
- Study startup managers who are accountable for activation timelines they cannot control.
- Site directors and PIs who need inspection readiness without the inspection prep panic.
- Clinical operations leaders who want measurable efficiency gains across their portfolio.

Agents

Six purpose-built AI agents for the operational problems that cost your team the most time, money, and risk. Each agent is designed for a specific workflow, validated against clinical trial standards, and ready to use in minutes.

Protocol-to-Site-Operations Generator

Turn a protocol into a complete site operational package. In 60 seconds.

Every new study starts the same way. A coordinator receives a 200-page protocol and spends 40 to 80 hours manually extracting the information they need to run the study: visit schedules, document checklists, training requirements, source document worksheets. This process is repeated at every

single site, for every single study. Thousands of hours of duplicated effort across the industry every year.

The Protocol-to-Site-Operations Generator reads the protocol once and produces the complete operational package. Five deliverables, ready for review, in the time it takes to pour a cup of coffee.

Problem: Coordinators spend 40-80 hours per study manually interpreting protocols into operational documents. This work is repeated identically at every activated site.

Built for: Clinical research coordinators, study startup managers, CRAs

How it works: Upload a protocol PDF. Flo State reads every section and generates five operational deliverables: (1) visit schedule table with time windows, procedures, and estimated visit duration, (2) document checklist mapped to TMF Reference Model zones, (3) role-based training requirements matrix for PI, sub-I, CRC, and support staff, (4) source document worksheet outlining what source data each CRF field requires, (5) site burden estimate with hours-per-role breakdown and predicted activation timeline.

What you get: Five-tab output with complete visit schedule, 40-60 line document checklist, training matrix by role, source document worksheet, and burden dashboard. All downloadable.

Time saved: 40-80 hours per study, per site. Across a 50-site study, that is 2,000-4,000 hours returned to the system.

Example: A new CRC joins Metro Clinical Research Center and is assigned her first oncology study. Instead of spending her first two weeks manually building study binders, she uploads the protocol to the Generator. Within a minute she has her visit schedule, knows exactly which regulatory documents to collect, sees that she and the PI both need RECIST training, and has a source document worksheet ready for the first patient visit. She spends her first two weeks learning the science, not the paperwork.

▸ DEMO WORKFLOW: Protocol-to-Site-Operations Generator

Route: /agents/protocol-to-site-operations

Step	Action	What Happens
1	From StoreFront, click Protocol-to-Site-Operations Generator	Opens detail page with description
2	Click Activate Agent	Upload screen with file dropzone and “Process Protocol” button
3	(Optional) Drop a protocol PDF	File name shown, but processing uses embedded DERM-110 data
4	Click Process Protocol	3s processing animation
5	Results load with 5 tabs	Visit Schedule shown first
6	Visit Schedule tab	6-visit timeline (Screening → Baseline → Day 7/14/21/28) with procedures, duration, notes. “Add to Google Calendar” opens 6 pre-filled calendar events
7	Click Export Operations Bundle	Downloads: DERM110_Operations_Bundle.csv
8	Click Regulatory Checklist tab	Essential docs list (IRB Approval, 1572, Delegation Log, Lab Certs) with status badges and responsible parties
9	Click Training Matrix tab	Role-based training (PI, Sub-I, CRC, Research Nurse, Pharmacist, Lab Tech) with completion % per role
10	Click Source Documents tab	8 source doc templates (Vital Signs, ECG, PASI/IGA Scoring, AE Log, etc.) each with Download button

11	Click any source doc Download	Downloads individual .txt file (e.g., Vital_Signs_ECG_Worksheet.txt)
12	Click Reference Card tab	Single-page study summary card with protocol highlights, key contacts, visit window rules

📁 Input: None — embedded mock data (DERM-110 Plaque Psoriasis trial)
📄 Output 1: DERM110_Operations_Bundle.csv (from Export Operations Bundle)
📄 Output 2: {DocTitle}.txt per source doc template
🔗 Output 3: 6 pre-filled Google Calendar event links (opens browser)

eISF-eTMF Reconciliation Bot

Find every document mismatch between site and sponsor. In 10 seconds.

CRAs spend 3 to 7 hours per site, per visit, cross-checking site documents against the sponsor's Trial Master File. Version mismatches, missing filings, expired credentials, naming inconsistencies. These are not clinical problems. They are filing problems. And when they are not caught, they become inspection findings.

The Reconciliation Bot compares every document in the site's eISF against the sponsor's eTMF and generates a severity-rated discrepancy report. Critical findings surface first. Minor discrepancies are cataloged for follow-up. CRAs walk into monitoring visits already knowing where the issues are.

Problem: Document mismatches between site eISF and sponsor eTMF are the most common monitoring finding and the leading cause of inspection observations. Manual reconciliation takes hours and still misses discrepancies.

Built for: CRAs, TMF managers, quality assurance teams, site regulatory specialists

How it works: Flo State compares document metadata across both systems: version numbers, file names, signatures, filing locations, expiration dates, and upload timestamps. Nine reconciliation rules run simultaneously, covering version mismatches, missing documents, naming inconsistencies, filing location errors, expired documents, signature gaps, date discrepancies, and stale uploads. Each finding is severity-rated: Critical, Major, or Minor.

What you get: Split-view reconciliation report showing matched and mismatched documents side by side. Each discrepancy includes the rule that triggered it, severity rating, specific description, and recommended corrective action.

Time saved: 3-7 hours per site per monitoring visit. For a CRA monitoring 10 sites monthly, that is 30-70 hours returned per month.

Example: A CRA prepares for a routine monitoring visit at a community hospital site. She runs the Reconciliation Bot before driving to the site. In 10 seconds, she sees three findings: the ICF in the eISF is version 3.0 but the eTMF still has version 2.0 (Critical), a sub-investigator's financial disclosure is missing from the eISF (Critical), and a training certificate is filed in Zone 03 instead of Zone 05 (Major). She contacts the site before the visit to request the ICF update, arrives prepared, and completes the visit in half the usual time.

▸ DEMO WORKFLOW: eISF-eTMF Reconciliation Bot

Route: /agents/eisf-etmf-reconciliation-bot

Step	Action	What Happens

1	From StoreFront → Agents → click eISF-eTMF Reconciliation Bot	Opens detail page with description, “Built For,” “Time Saved”
2	Click Activate Agent	Upload screen with two file dropzones (eISF export + eTMF export) and Run Reconciliation button
3	(Optional) Drag/drop files into dropzones	Files accepted but reconciliation uses embedded JSON data regardless
4	Click Run Reconciliation	1.5s processing — cross-references all docs comparing versions, missing docs, expired docs, filing errors, naming, signatures
5	Report renders	Discrepancy table: Severity, Artifact, TMF Zone, eISF Value, eTMF Value, Action. Color-coded: 🔴 Critical, 🟡 Major, 🔵 Minor
6	Click Export Report	Downloads: eISF_eTMF_Discrepancy_Report.csv
7	Click Generate Resolution Plan	Filters to Critical + Major discrepancies, presents ordered action plan with pending badges
8	Click Execute All	Agentic execution loop — each action runs sequentially with real-time terminal log
9	All actions complete	“All Actions Complete” heading with full execution log visible

📁 Input: src/data/eisf_documents.json + src/data/etmf_documents.json (embedded)
📄 Output: eISF_eTMF_Discrepancy_Report.csv (from Export Report button)

Virtual SMO Layer

Your coordinator’s AI-powered morning briefing. Every study. Every task. Every alert.

China enrolls patients 2-5x faster than the rest of the world. One reason: embedded Site Management Organizations (SMOs) that handle administrative and operational burden so clinical staff can focus on patients. Western sites do not have SMOs. They have coordinators managing 7.6 studies simultaneously, with 20+ systems, 6+ logins, and no unified view of what matters today.

The Virtual SMO Layer is the digital version of an embedded operations team. It runs before your coordinator opens Florence in the morning, curating a prioritized view of everything that needs attention across every study: documents staged for today’s monitoring visit, credentials expiring this week, draft responses to sponsor queries ready for one-click approval, scheduling conflicts identified for tomorrow, and a weekly portfolio status dashboard. This is not a notification system. It is an operations layer that thinks ahead.

Problem: Coordinators manage 7+ studies across 20+ systems with no unified operational view. Critical tasks fall through cracks. Reactive work crowds out proactive work. Burnout drives 35-61% annual turnover.

Built for: Clinical research coordinators, site directors, research managers

How it works: Flo State aggregates data across all active studies for a specific coordinator: document status, staff credentials, study schedules, sponsor communications, and regulatory deadlines. It prioritizes tasks by urgency and impact, pre-stages documents for scheduled visits, drafts responses to pending queries, and identifies scheduling conflicts 24-48 hours in advance. The coordinator sees a single, curated dashboard every morning.

What you get: Morning priorities dashboard (5-7 ranked tasks across all studies), monitoring visit prep package with pre-staged documents and completeness indicators, regulatory alerts for expiring credentials and upcoming renewals, AI-drafted communications ready for review and one-click

approval, scheduling conflict detection with alternative time suggestions, and weekly study portfolio status with red/yellow/green indicators.

Time saved: 2-3 hours per coordinator per day. Across a 10-coordinator site running 25 studies, that is over 5,000 hours per year returned to patient-facing work.

Example: Maria Rodriguez, CRC at Metro Clinical Research Center, opens Flo State at 8:00 AM. Her dashboard shows: (1) monitoring visit at 10:30 for Study CNS-205, documents pre-staged and 100% complete; (2) Dr. Patel's GCP certification expires in 14 days, renewal link included; (3) draft response to Pfizer query about a protocol deviation in Study ONK-301, ready for review; (4) scheduling conflict detected for tomorrow: two monitoring visits from different sponsors overlap at 2:00 PM, with a suggested reschedule for one. Maria reviews and approves the query response, forwards the GCP renewal to Dr. Patel, accepts the reschedule, and walks into her first patient visit fully prepared. The work that used to take her entire morning happened before she finished her coffee.

▸ DEMO WORKFLOW: Virtual SMO Layer

Route: /agents/virtual-smo-layer

Step	Action	What Happens
1	From StoreFront, click Virtual SMO Layer	Opens detail page
2	Click Activate Agent	Morning Briefing phase opens immediately (no upload step)
3	Briefing shows	Personalized greeting: "Good morning, Maria" with site info, date, workload gauge. AI narrative summarizes active studies, pending tasks, credential alerts
4	Review Priority Tasks	Task cards with P1-Critical through P4-Routine badges, study name, deadline. Click ✓ to dismiss
5	Review Credential Alerts	Expiring certifications with countdown badges (e.g., "14 days"). Click to dismiss
6	Scroll through remaining sections	Monitoring Visit Prep, Drafted Communications, Scheduling Conflicts, Study Portfolio — all pre-populated
7	Click Export Portfolio	Downloads: SMO_Weekly_Portfolio_Status.csv
8	Click Generate Execution Plan	Plan phase — 6 AI-recommended actions with study context, each marked "pending"
9	Review 6 execution plan items	E.g., "Auto-reply to sponsor query (ONK-301)," "File amended delegation log (CARD-112)," "Pre-fill IRB continuing review (ENT-225)"
10	Click Execute All Actions	Agentic execution — each item runs sequentially with real-time terminal log (~3-4s each, ~20s total)
11	All 6 actions complete	"All Actions Complete" heading. Full terminal log visible. Each item shows green ✓ "complete" badge

📁 Input: src/data/coordinator_maria.json (CRC profile, tasks, alerts, studies — embedded)
📄 Output: SMO_Weekly_Portfolio_Status.csv (from Export Portfolio button)

Templates

Pre-configured AI workspaces with the right documents loaded, the right instructions written, and the right skills activated. Install one and start working. No setup. No AI expertise required.

CRA Monitoring Visit Prep

Walk into every visit knowing exactly what to look for.

CRAs prepare for monitoring visits by reviewing site documents, checking regulatory status, and building visit-specific checklists. The preparation is repetitive, time-consuming, and prone to gaps. And when a CRA misses something during prep, it becomes a missed finding during the visit.

This template creates a monitoring-focused AI workspace that understands ICH E6(R3) monitoring requirements, TransCelerate risk-based monitoring principles, and the 50 most common FDA/EMA inspection findings. Use it to generate visit-specific checklists, identify priority review areas when time is limited, and draft follow-up letters.

Built for: CRAs, lead CRAs, monitoring managers

What's included: 4 knowledge documents: ICH E6(R3) Section 5.18 (Monitoring), TransCelerate RBM Guidelines, TMF Completeness Criteria by Visit Type, Common Monitoring Findings Database (top 50 FDA/EMA findings). 2 pre-enabled skills: Document Comparator, Regulatory Calendar.

What you can ask: 'Generate a monitoring checklist for a routine visit at an oncology site.' / 'I only have 3 hours at this site. What should I prioritize?' / 'Draft a follow-up letter noting 2 minor documentation findings.' / 'What are the most common inspection findings for informed consent?'

Why this exists: CRAs monitor 8-12 sites monthly. Standardized, intelligent visit preparation reduces missed findings and shortens visit time.

▸ DEMO WORKFLOW: CRA Monitoring Visit Prep

Route: /templates/cra-monitoring-visit-prep

Step	Action	What Happens
1	From StoreFront, click Templates tab	Shows template cards
2	Click CRA Monitoring Visit Prep card	Opens detail page with description
3	Click Install Template	3-stage install animation (~5s): Scaffolding → Ingesting → Activating
4	Animation completes → workspace opens	Visit info banner + 6-section checklist loads
5	Click checklist items to check them off	Items show strikethrough, progress bar updates, FLAGGED/PASS badges disappear
6	Collapse/expand sections via chevron	Sections toggle, showing per-section progress (e.g., 3/7)
7	Check all 25 items	Green "Pre-Visit Checklist Complete" banner appears with "Generate Visit Report" button
8	Click Export Report (top right)	(Demo only — no file output)

📁 Input: None — all mock data embedded (ONK-301 study, Site 4521, CRA Jennifer Walsh)
📄 Output: None — Export Report is demo-only (no file download wired)

Skills

Modular AI capabilities that snap into any agent or template. Enable a skill once and it works everywhere. Composable, reusable, and always improving.

Protocol Reader

Turns any protocol PDF into structured, queryable data.

Clinical trial protocols are dense, unstructured documents. The information your team needs (visit schedules, eligibility criteria, endpoint definitions, study design parameters) is scattered across 100-300 pages of text. Every time someone needs a specific detail, they search, scroll, and hope they find the right section.

Protocol Reader extracts structured data from any protocol PDF: visits, procedures per visit, eligibility criteria (inclusion and exclusion), primary and secondary endpoints, study arms, and design parameters. Once extracted, the data is available to every agent and template in your workspace. Upload once, reference everywhere.

Built for: All clinical trial roles

Accepts: Protocol PDF files (.pdf, text-extractable)

Produces: Structured JSON: visits (with procedures), eligibility criteria, endpoints, study design, metadata

Works with: Protocol Complexity Scorer, Protocol-to-Ops Generator, Amendment Impact Assessment template, Feasibility Response template

▸ DEMO WORKFLOW: Protocol Reader

Route: /skills/protocol-reader

Step	Action	What Happens
1	From StoreFront → Skills → click card	Opens detail page
2	Click Activate Skill	Upload screen appears with two options
3	Click Load Sample Protocol	2.5s loading animation: “Extracting Protocol Data...” with 5 parsing steps
4	Extraction completes → Study Overview tab	PHOENIX-301 metadata: Sponsor, Phase III, 450 subjects, 1:1 randomization, 2 study arms
5	Click Schedule of Events tab	24-visit × 16-procedure matrix with checkmarks (incl. PK Sampling note “Cohort 1 only”)
6	Click Eligibility tab	6 Inclusion + 5 Exclusion criteria with plain-English translations and “Restrictive” / “Screen Failure Risk” badges
7	Click Endpoints tab	5 endpoints (1 Primary PFS, 3 Secondary, 1 Exploratory) with definitions and statistical methods
8	Click Design Parameters tab	4 cards: Sample Size, Interim Analyses, Stratification, DSMB
9	Click Complexity tab	4 KPIs (24 visits, 16 procedures, 263/patient, ~3.1M data points) + benchmark comparison table
10	Click Export All (top right)	(Demo only — no file output)

> 📁 Input: “Upload Protocol PDF” button exists but is demo-only. “Load Sample Protocol” loads embedded PHOENIX-301 data

📄 Output: None — Export All is demo-only (no file download wired)

Audit Trail Generator

GxP-ready documentation of every AI-assisted action.

Clinical trials operate under Good Clinical Practice (GCP) regulations that require complete documentation of every action that affects trial data, patient safety, or regulatory compliance. When AI assists with document preparation, protocol interpretation, or quality review, regulators need to know what the AI did, what data it used, and what the human decided.

Audit Trail Generator runs in the background during any Flo State session, logging every AI-assisted action: inputs provided, outputs generated, user review decisions (approved, edited, rejected), and timestamps. The output is a compliance-ready audit trail that can be included in the Trial Master File or presented during inspections.

Built for: Quality managers, regulatory specialists, site directors, anyone working in a GxP environment

Accepts: Session activity log (automatic)

Produces: Formatted audit trail (downloadable as PDF or JSON) with timestamps, action descriptions, input/output hashes, and user decisions

Works with: All agents and templates. Always available.

▸ DEMO WORKFLOW: Audit Trail Generator

Route: /skills/audit-trail-generator

Step	Action	What Happens
1	From StoreFront → Skills → click card	Opens detail page
2	Click Activate Skill	Session log opens immediately (no upload step)
3	Session Log view loads	12 timestamped audit entries spanning a full day, each showing user, agent, action type, and Approved/Edited/Rejected badge
4	Click any entry to expand	Shows full Input Summary, Output Summary, Study, Document, Confidence %, and Decision
5	Click GxP filter buttons (e.g., GxP-critical)	Filters to only GxP-critical entries (5 entries)
6	Click Session Summary tab	4 KPI cards (12 Total, 9 Approved 75%, 2 Edited 17%, 1 Rejected 8%) + session details (7h 13m, 2 users, 7 agents, 9 docs, 3 studies)
7	Click Compliance Dashboard tab	Human-in-the-loop metrics (25% override, 75% AI accuracy, 92% avg confidence) + ALCOA+ checklist (9/9 green) + “Regulatory Readiness: 100%” banner
8	Click Export for TMF (top right)	(Demo only — no file output)
9	Click Inspector View button	(Demo only — styled but no action)

📁 Input: None — 12 mock audit entries embedded (references Virtual SMO, Reconciliation Bot, Protocol Reader, Regulatory Factory, and other agents)
📄 Output: None — Export for TMF and Inspector View are demo-only

Cash Flo

The company-wide dashboard that translates every Flo State action into dollars, days, and decisions.

Every agent, template, and skill in Flo State generates measurable operational value. Cash Flo is where that value becomes visible. It is the real-time executive dashboard that connects AI activity to financial outcomes: value reclaimed, activation days accelerated, findings auto-resolved, hours returned to clinical work. Cash Flo turns Flo State from a productivity tool into a boardroom-ready investment thesis.

Cash Flo is designed for three audiences simultaneously: clinical operations leaders who need to see operational velocity improvements, finance teams and PE boards who need EBITDA impact and margin expansion data, and technology leaders who need to justify AI token investment with concrete ROI. Six tabs, one story: AI is working, and here is the proof.

▸ DEMO WORKFLOW: Cash Flo

Route: Top navigation from StoreFront

Step	Action	What Happens
1	From StoreFront, click Cash Flo in top nav	Dashboard opens on Overview tier
2	Overview — note 4 hero KPIs	$12.4M Value Reclaimed (+312% YoY), 14,500 Days Activation Accelerated, 28,450 Findings Auto-Resolved, 142k Hours Reclaimed
3	Observe Token Usage vs Value area chart	Dual-axis: token consumption (bars) vs savings (purple line) over weekdays
4	Observe Top Agent Activity bar chart	Virtual SMO leads with 34,200 calls, followed by CRA Visit Prep, Protocol Scorer, etc.
5	Observe Top Operators and Top Builders tables	Operators ranked by value generated; Builders ranked by agents/templates/skills created
6	Toggle to This Month	All KPIs update: $48.2M value, 60,200 activation days, charts shift to week-level granularity
7	Toggle to This Year	KPIs jump to $412.5M, 725k days, charts show quarterly data
8	Click Operational Velocity tab	
9	Observe Document Cycle Time	Line chart: 30-day rolling avg, AI-Assisted (purple) vs Non-AI (gray). Callout: “3.2× Speed Multiplier.” KPIs: 2.1 day avg, 87% first-pass, 45% reduction
10	Observe eISF Activation	Area chart showing activation phases. KPIs: 342 days saved, 47 sites, 7.2 avg days
11	Observe TMF Inspection Readiness	Readiness bars per study (ONK-301 94%, CARD-112 88%), green/amber/red. KPI: 28,450 findings auto-resolved
12	Click Site Performance tab	
13	Observe Engagement vs Enrollment scatter	Sites plotted by engagement (x) vs enrollment (y), trend line showing correlation
14	Observe Adoption Funnel	Vertical funnel: Licensed (8,500) → Active (6,800) → First AI Use (5,100) → Regular (3,400) → Power Users (1,200)
15	Click Financial Impact tab	
16	Observe Revenue Conversion	Stacked area: backlog → ARR → AI-accelerated overlay across quarters

17	Observe EBITDA Margin Bridge	Waterfall: 24.2% → +0.8% Labor → +0.5% Cost Avoidance → +0.3% Rework → +1.1% CRA Capacity → 26.9%. Callout: “270 bps margin improvement”
18	Observe Churn & Retention	3-line chart: High Adoption (98%), Medium (92%), Low (84%). NRR gauge at 124%
19	Click Predictive Intel tab	
20	Observe Risk Radar	Radar chart: Enrollment Failure, Regulatory Lapse, Quality Degradation, Site Disengagement, Amendment Disruption
21	Observe Sites Needing Attention	Flagged sites table with health score, issue, trend indicator (↓ declining, → stable)
22	Observe Prediction Accuracy	Bar chart: monthly accuracy 82-91%, callout: “87% accuracy over last 90 days”
23	Click Token Economics tab	
24	Observe Agent Efficiency Ranking	All agents ranked by tokens consumed, value generated, value per token, usage count
25	Observe Unit Economics KPIs	4 cards: $0.0034 avg cost/token, $2.84 value/token, 835× ROI ratio, 12.4M hours equivalent

📁 Input: None — all data aggregated from Flo State platform activity
📄 Output: None — dashboard is view-only (no export wired in demo)

Cash Flo — Key Demo Talking Points by Audience

Tab	Killer Stat	Audience
Overview	$12.4M Value Reclaimed	Everyone
Op. Velocity	3.2× Speed Multiplier on document cycle time	ClinOps
Site Performance	Sites using ≥3 tools enroll 2.1× faster	Site Networks
Financial Impact	270 bps EBITDA margin improvement	PE Board / CFO
Predictive Intel	87% prediction accuracy, $4.2M avoided	Risk Mgmt
Token Economics	835× ROI ratio on AI token investment	CTO / IT

Recommended 5-7 Minute Hackathon Demo Order

Four capabilities, one narrative arc: from “what it does” to “what it’s worth.”

#	Capability	Time	Why This One
1	**Protocol-to-Site-Ops Generator**	~90s	Most visceral “wow” moment. Upload, 3 seconds, five tabs of output. Audience immediately gets what Flo State does.
2	**eISF-eTMF Reconciliation Bot**	~90s	Shows the agentic execution loop. Run reconciliation, show severity report, hit Execute All. Terminal log proves AI is doing work, not just showing work.
3	**Protocol Reader**	~45s	Quick hit: activate, load sample, flip through Schedule of Events + Complexity. Shows Skills layer and how structured data feeds agents.

4	**Cash Flo (Closer)**	~2-3m	Start Overview hero KPIs, toggle to This Year, walk Financial Impact (EBITDA waterfall), finish Token Economics (835× ROI). Ties everything to dollars.

How We Built It

We built this in 48 hours using a multi-model AI stack: Claude Opus 4.6 for deep protocol analysis and structured extraction, Gemini 3.1 Pro for rapid code generation and iteration, and Google Antigravity as the agentic coding orchestrator that wired it all together.

The architecture uses a hexagonal adapter pattern so every agent can run live against the LLM or fall back to domain-expert cached outputs — the demo works perfectly offline.

The data layer is a clinical trial knowledge graph mapping protocols, TMF documents, site profiles, and coordinator workflows using the DIA TMF Reference Model ontology.

Agents run in an agentic execution loop — they analyze, generate a plan, then autonomously execute multi-step workflows with real-time streaming logs.

The frontend is React + TypeScript + Recharts with Playwright E2E tests validating every workflow.

Getting Started

Flo State is available inside Florence. No separate login. No additional software. No AI training required.

Open Flo State. Browse agents, templates, and skills. Activate what you need. Start working.

Every capability described in this catalog is purpose-built for clinical trial operations, validated against industry standards, and designed to keep your team in control while removing the manual work that slows them down.

That is what ClinOps can look like in the Flo State.